

# Python's Superpower for Numbers

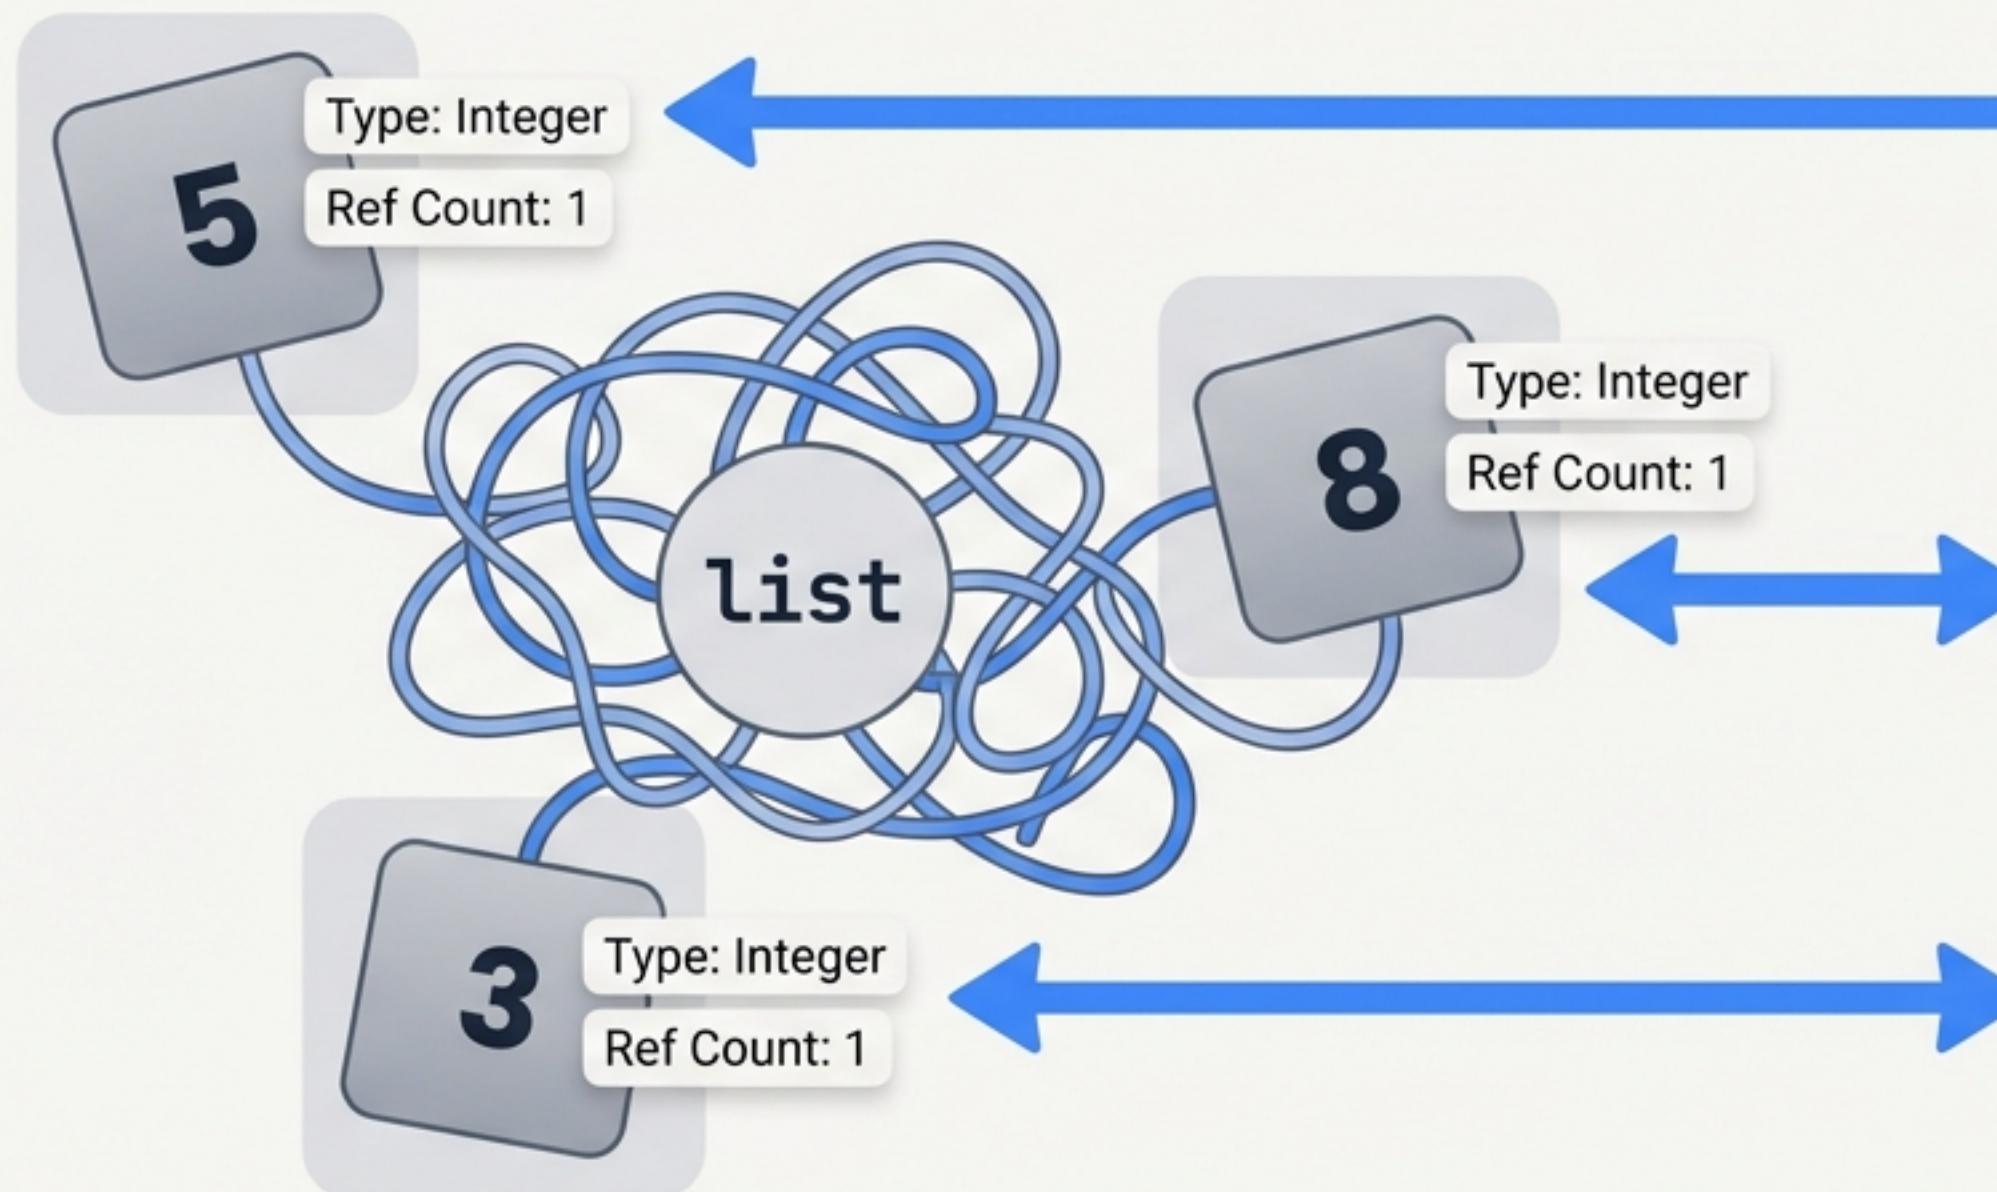
A Beginner's Guide to Mastering NumPy



# First, Why Do We Even Need NumPy?

Standard Python is fantastic, but it wasn't built for high-speed, large-scale data crunching. Its core data types, like lists, have a fundamental problem.

## Python List: The Jumbled Toy Box



## The Problem

### Memory Overhead

Each number isn't just a number; it's a full-blown Python object with extra metadata, wasting space.

### Pointer Chase

The CPU has to follow each tangled pointer to find the next number in memory. This is slow and inefficient.

### Type Checking

For every single operation (like adding two numbers), Python has to check the type of each box.



# The Real Cost: Python Loops are Incredibly Inefficient

```
# Adding two lists in Python
c = []
for i in range(len(a)):
    c.append(a[i] + b[i])
```

This happens for millions of numbers. It's like gift-wrapping and unwrapping every single number for every single calculation.

What Python does for *\*every single number\**:

- 
- Follow a pointer to a box.
  - "Is this an integer?"
  - Extract the raw number.
  - Follow another pointer.
  - "Is this also an integer?"
  - Extract the raw number.
  - Perform the actual addition.
  - Create a brand new Python object (a new box) for the result.
  - Add a pointer to this new box to list `c`.



# NumPy to the Rescue!

NumPy solves Python's speed problem by going back to first principles with a new data type: the `ndarray`.

## Solution A: Contiguous Memory



- ✓ **A solid block:** A NumPy array is a continuous block of memory.
- ✓ **No 'boxes':** It stores the raw values directly, with no pointers or metadata overhead.
- ✓ **Homogeneous Data:** All elements must be the same type (e.g., all `int64`). This eliminates the need for type-checking.



### The Godfather of NumPy

Created by Travis Oliphant in 2005 to unify two competing libraries (Numeric and Numarray) and give Python a single, powerful tool for scientific computing.

# The Secret Weapon: Vectorization (SIMD)

Because the data is in one continuous block and the type is known, NumPy can use special, low-level CPU instructions called **SIMD (Single Instruction, Multiple Data)**. Instead of adding numbers one by one in a loop, the CPU can load a chunk (a “vector”) of 4 or 8 numbers at once and perform the operation in a single clock cycle.

**Python Loop**



Left Code Block (Python)

```
# Inefficient Loop
for i in range(len(a)):
    c.append(a[i] + b[i])
```

**NumPy Vectorization**



Right Code Block (NumPy)

```
# Vectorized Operation
c = a + b
```



This one line triggers the super-fast vectorized operation!

# 50-100X FASTER



VS



**Pure Python Loop**

**NumPy Vectorized Operation**

This is why NumPy is the foundation for almost every data science, machine learning, and scientific package in Python, including Pandas, Matplotlib, and Scikit-learn.

# Level 1: Creating Your First Arrays

## Step 1: The Import

```
import numpy as np
```



First, we summon our hero and give it a short name, `np`.

## Step 2: Creating Arrays from Lists

### Example 1: 1D Array (Vector)

```
a = np.array([1, 2, 3])
```

1	2	3
---	---	---

### Example 2: 2D Array (Matrix)

```
b = np.array([[1, 2, 3],  
              [4, 5, 6]])
```

1	2	3
4	5	6

### Example 3: 3D Array (Tensor)

```
c = np.array([[[1,2], [3,4]],  
              [[5,6], [7,8]]])
```



Notice the square brackets!  
A 3D array is just a list of  
2D arrays.

5	6	
1	2	
7		

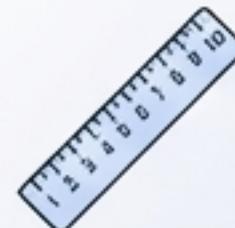
5	6	
1	2	
3	4	

# The Array Creation Toolkit

Create arrays from scratch without needing a Python list first.

## `np.arange(start, stop)`

Create a sequence of numbers, just like Python's `range()`.

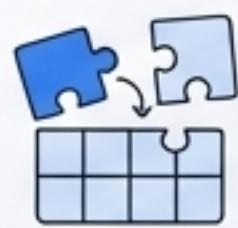


```
np.arange(1, 11)
```

```
[1 2 3 4 5 6 7 8 9 10]
```

## `.reshape(rows, cols)`

The perfect partner for `arange`! Molds your 1D array into a new shape.



```
np.arange(1, 13).reshape(3, 4)
```

```
[[1 2 3 4]
 [5 6 7 8]
 [9 10 11 12]]
```

## `np.zeros((rows, cols))`

You guessed it... an array full of 0s.



```
np.zeros((2, 3))
```

```
[[0. 0. 0.]
 [0. 0. 0.]]
```

## `np.linspace(start, end, num)`

Get `num` evenly spaced points between a `start` and `end`.

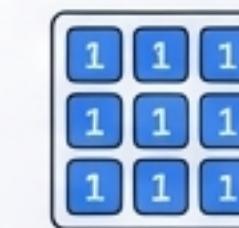


```
np.linspace(0, 10, 5)
```

```
[ 0.  2.5  5.  7.5  10.]
```

## `np.ones((rows, cols))`

Quickly make an array full of 1s. (Useful for initialization).



```
np.ones((2, 3))
```

```
[[1. 1. 1.]
 [1. 1. 1.]]
```

## `np.identity(n)`

Create an `n × n` identity matrix in one step.



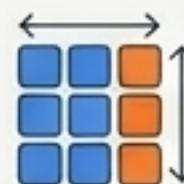
```
np.identity(3)
```

```
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
```

# Level 2: Inspecting Your Array

What's inside? Let's find out.

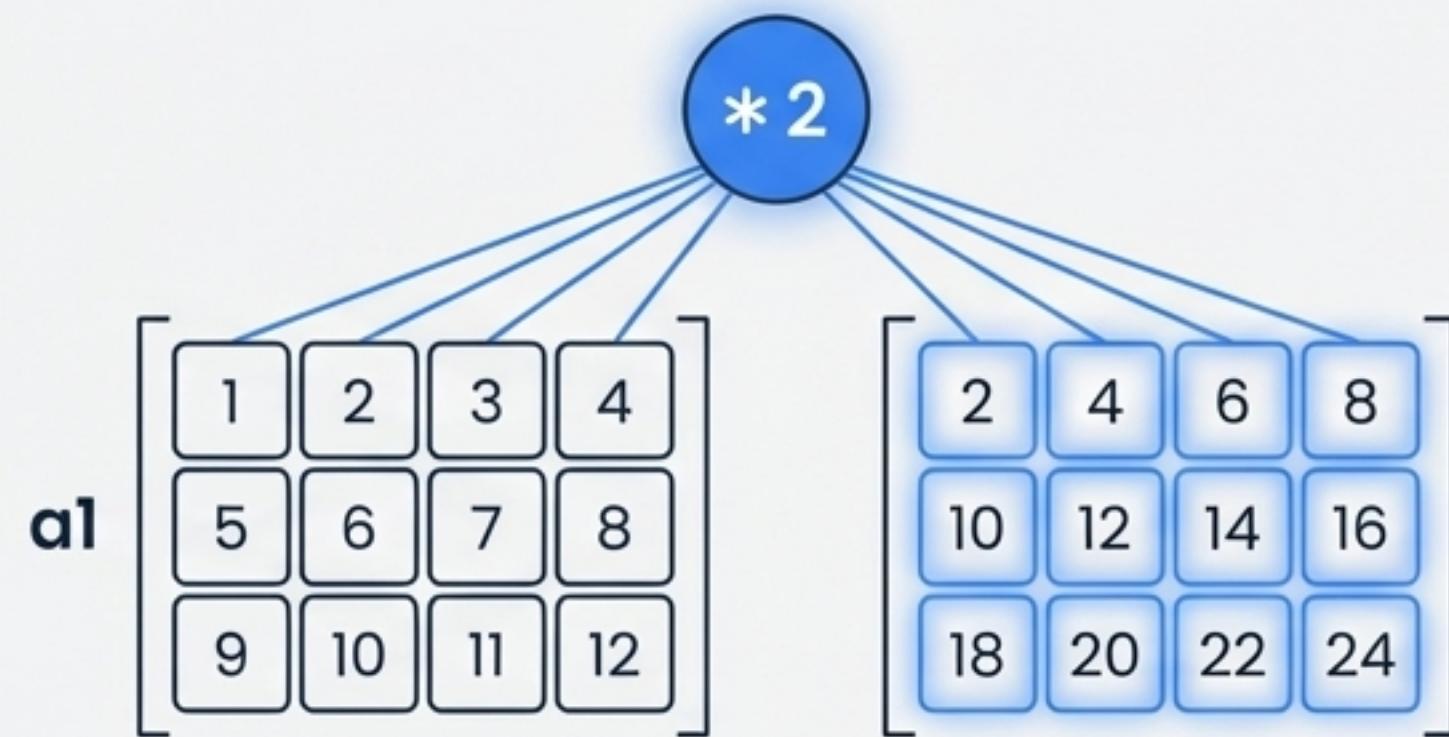
```
a = np.arange(12).reshape(3, 4)
# Our array for this section:
# [[ 0  1  2  3]
# [ 4  5  6  7]
# [ 8  9 10 11]]
```

Attribute	Output	What it means
a.ndim	2	Number of dimensions. It's a 2D matrix. 
a.shape	(3, 4)	The size of each dimension. 3 rows, 4 columns. 
a.size	12	Total number of elements in the array ( $3 * 4$ ). 
a.dtype	int64	The data type of the elements. Here, 64-bit integers. 
a.itemsize	8	How much memory (in bytes) each element takes up. 

# Operations on Entire Arrays, No Loops Needed.

## Part 1: Scalar Operations (Array + Number)

The operation is "broadcast" to every element in the array.



$a1 * 2$  (Multiplication)

```
a1 * 2
```

```
array([[ 2,  4,  6,  8],
       [10, 12, 14, 16],
       [18, 20, 22, 24]])
```

$a1 + 100$  (Addition)

```
a1 + 100
```

```
array([[101, 102, 103, 104],
       [105, 106, 107, 108],
       [109, 110, 111, 112]])
```

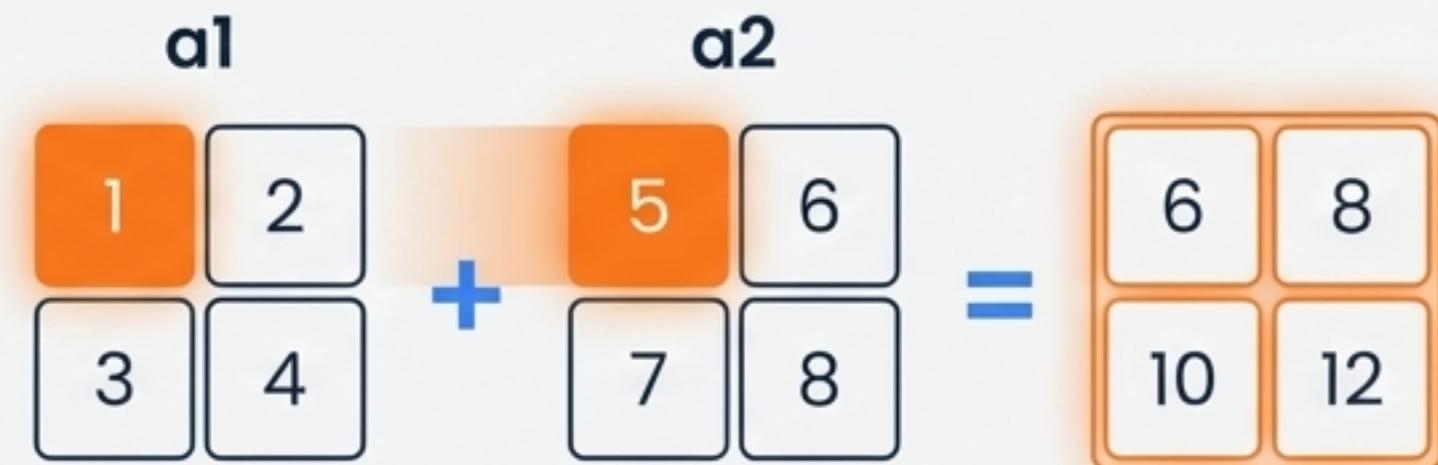
$a1 > 5$  (Relational)

```
a1 > 5
```

```
array([[False, False, False, False],
       [False, True, True, True],
       [True, True, True, True]])
```

## Part 2: Vector Operations (Array + Array)

The operation is performed element-wise between the two arrays.



$c = a1 + a2$

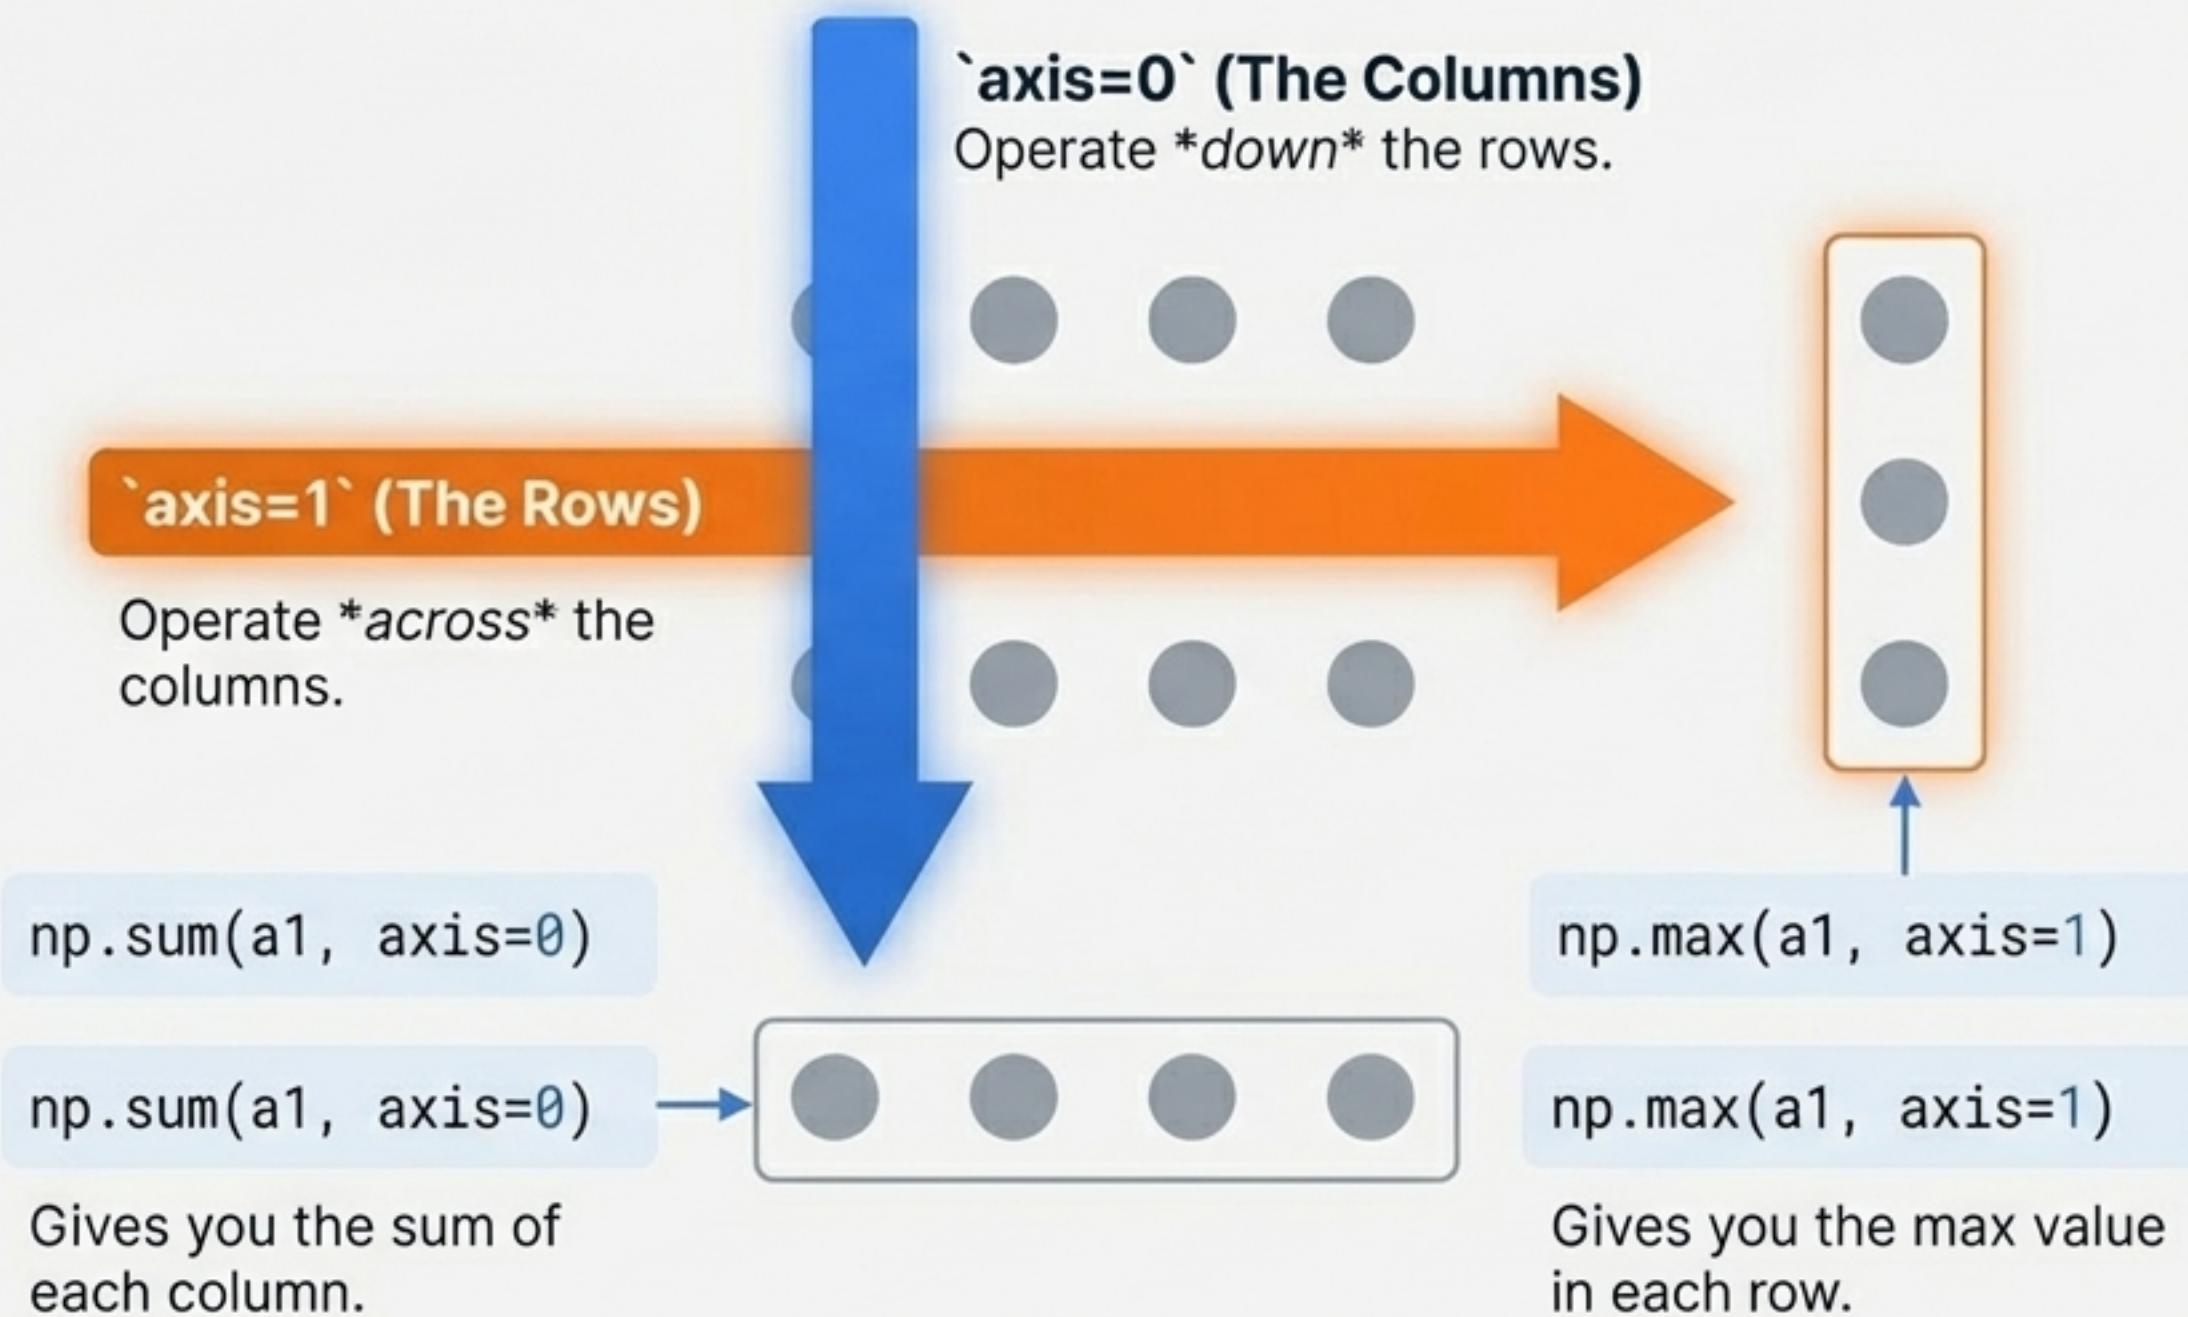
```
array([[ 6,  8],
       [10, 12]])
```

# The NumPy Function Toolkit

## Quick Stats & Aggregations

- `np.max(a1)`  
Find the max value in the entire array.
- `np.min(a1)`  
Find the minimum value.
- `np.sum(a1)`  
Sum all elements.
- `np.mean(a1)`  
Calculate the mean.
- `np.std(a1)`  
Calculate the standard deviation.

## The Most Important Parameter: `axis`



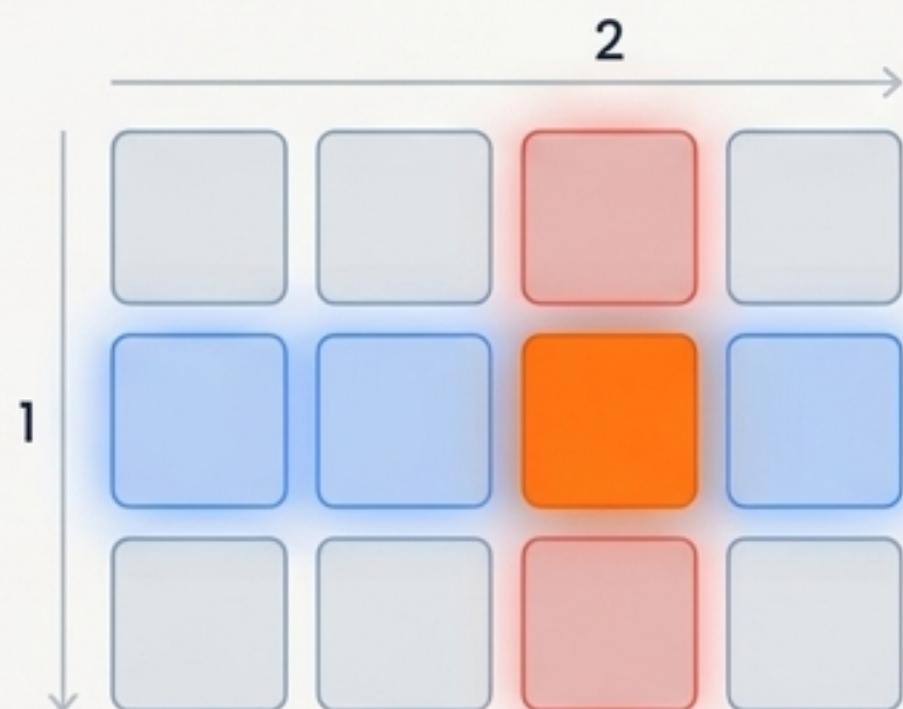
# Part 1: Indexing (Grabbing a Single Item)

1D Array



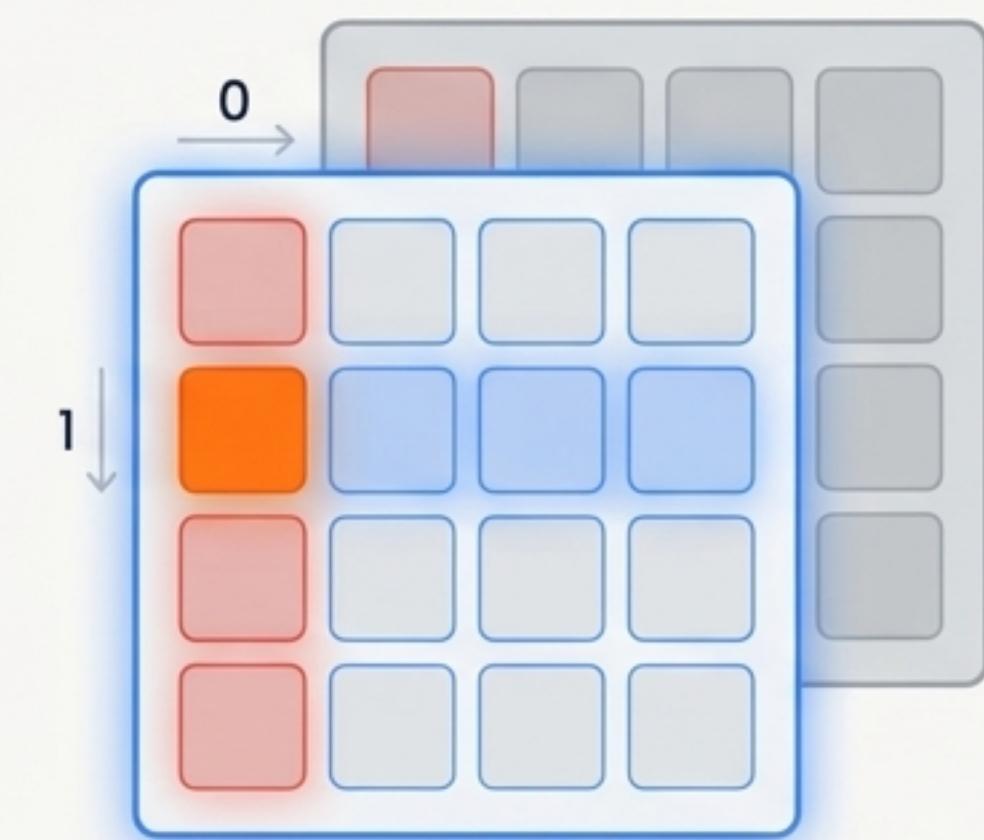
2D Array: `[row, column]`

a2[1, 2]



3D Array: `'[block, row, column]'`

a3[0, 1, 0]



# Part 2: Slicing (Grabbing Chunks of Items)

Use the colon ‘:’ which means “from the start to the end”.

	0	1	2	3	4
0	■	■	■	■	■
1	■	■	■	■	■
2	■	■	■	■	■
3	■	■	■	■	■

a2

## Get a full row

```
a2[1, :]
```

Get me Row 1, and all (‘:’) of its columns.

	0	1	2	3	4
0	■	■	■	■	■
1	■	■	■	■	■
2	■	■	■	■	■
3	■	■	■	■	■

## Get a full column

```
a2[:, 2]
```

Get me all (‘:’) the rows, but only from Column 2.

	0	1	2	3	4
0	■	■	■	■	■
1	■	■	■	■	■
2	■	■	■	■	■
3	■	■	■	■	■

## Get a sub-matrix

```
a2[1:3, 2:4]
```

Rows 1 up to (not including) 3. Columns 2 up to (not including) 4.

	0	1	2	3	4
0	■	■	■	■	■
1	■	■	■	■	■
2	■	■	■	■	■
3	■	■	■	■	■

## Using Steps

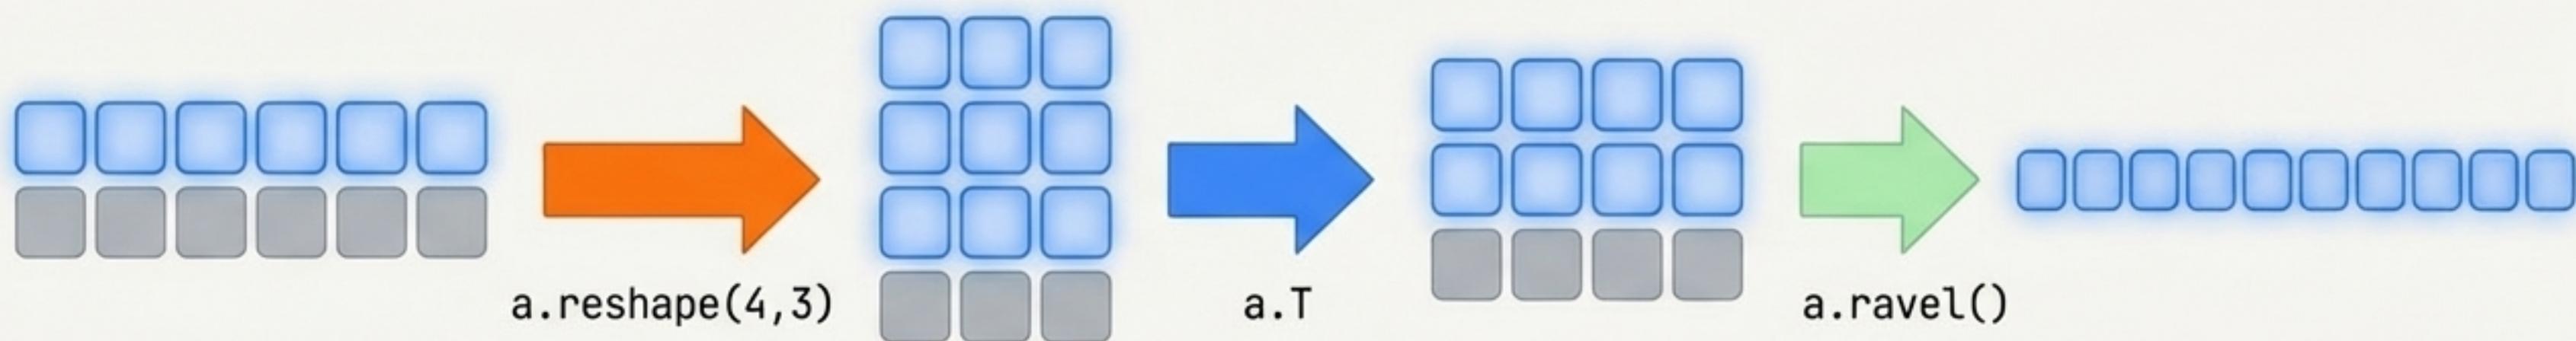
```
a2[::2, ::2]
```

From start to end, give me every 2nd row and every 2nd column.

	0	1	2	3	4
0	✓				✓
1					
2	✓			✓	
3					

# Reshaping Your World

Sometimes you need to change the layout of your data without changing the data itself.



Initial Array: (2, 6)

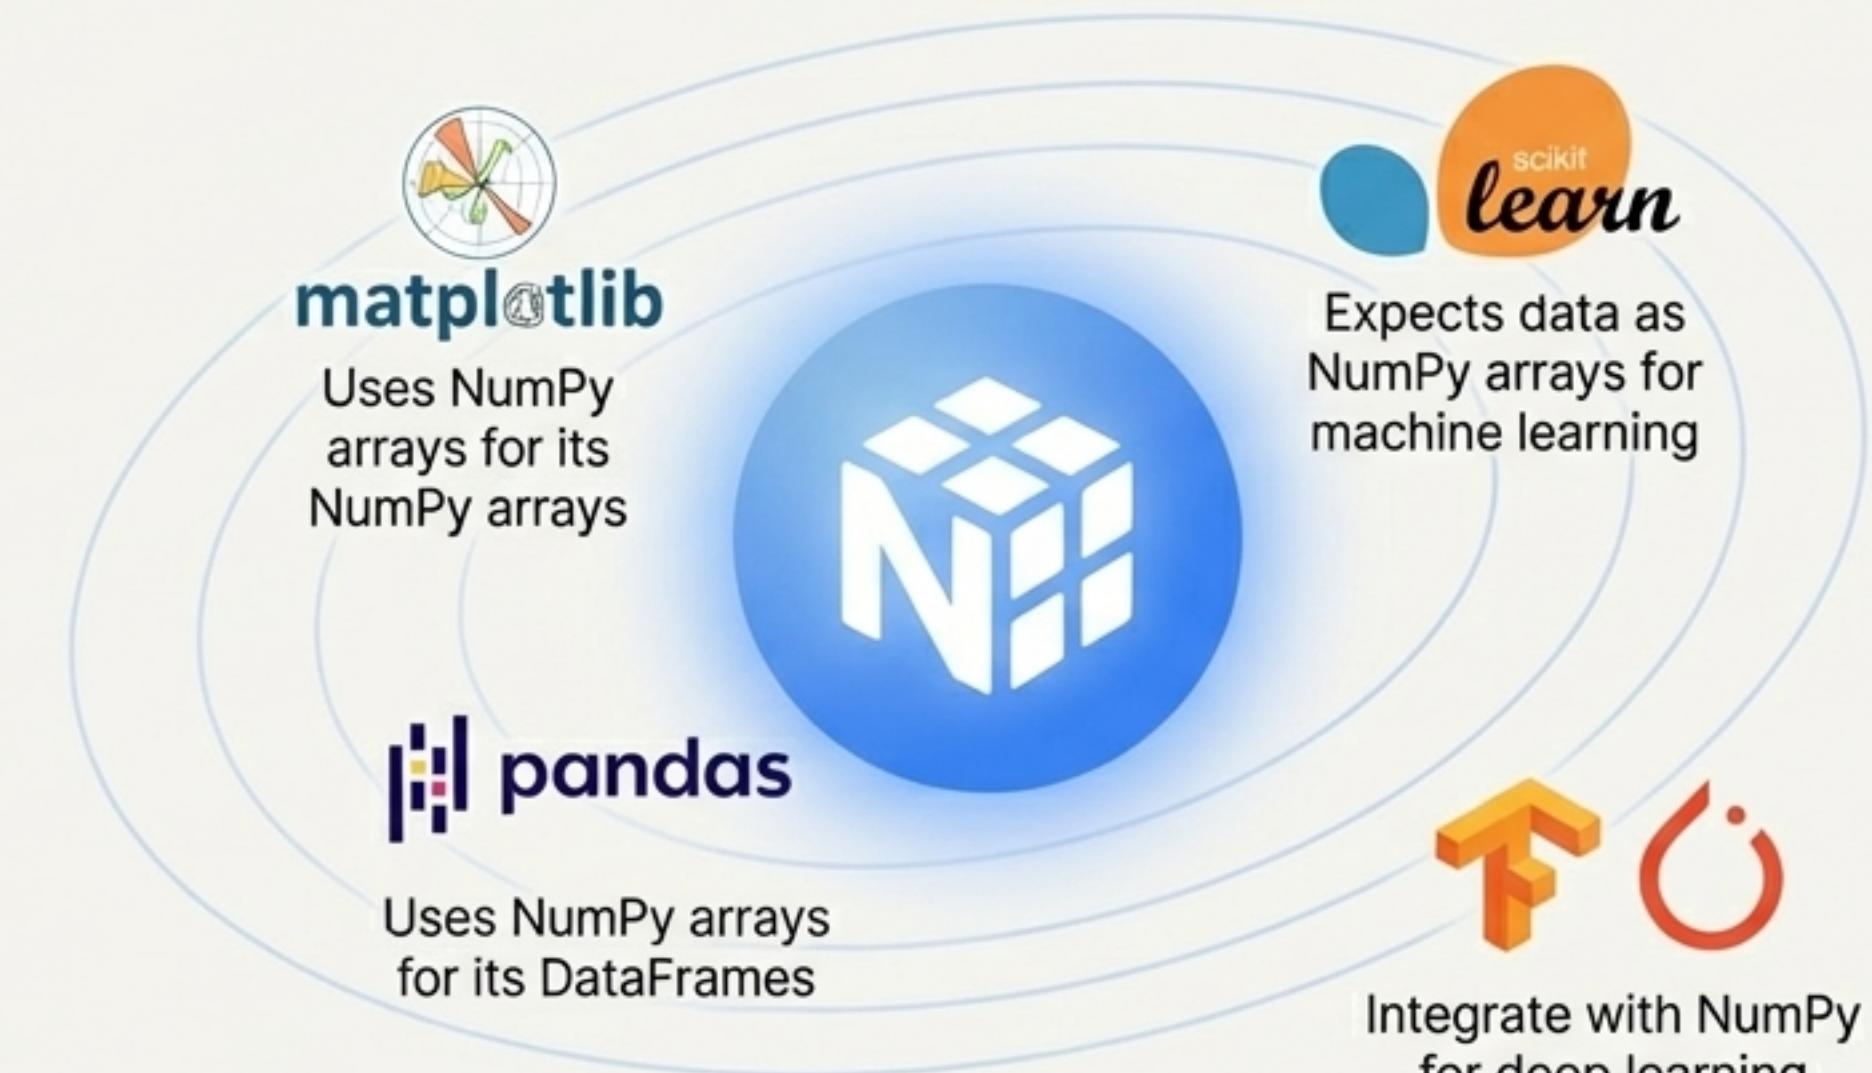
**Reshape:** Change the dimensions of an array, like a 2x6 into a 4x3. The total number of elements (**size**) must remain the same!

**Transpose:** Flip the rows and columns. A 4x3 becomes a 3x4.

**Ravel:** Flatten any array, no matter its dimension, into a single 1D row.

# The Foundation of Python's Data Ecosystem

- NumPy solves Python's speed problem with its `ndarray` object, using contiguous memory and vectorization.
- It gives you a powerful, loop-free syntax for mathematical and logical operations on entire datasets.
- It provides the fundamental array object that nearly every other major data library is built upon.



**Master NumPy, and you've mastered the building block  
of modern data science in Python.**